



Linear Regression

Qinliang Su (苏勤亮)

Sun Yat-sen University

suqliang@mail.sysu.edu.cn

Outline

- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization

Introduction

- What is regression?

Based on the given features, predict the values of interested variables

- Example: House price prediction

Features				Interest variable
Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
....

Features: 1) size, 2) # bedrooms, 3) # floors, 4) # years

- Mathematically, regression aims at building a function $f(\cdot)$ to model the relation between input data x and supervised value y

$$\hat{y} = f(x_1, x_2, x_3, x_4)$$

- **Linear** regression

Restricting the function $f(\cdot)$ to be of **linear form**, i.e.,

$$f(x_1, x_2 \dots x_m) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_m x_m$$

- w_k : model parameters
- m : number of features

- Objective

Find a set of parameters $\{w_k\}_{k=1}^m$ so that the prediction

$$\hat{y} = f(x_1, x_2, \dots, x_m)$$

is **as close as possible** to the true y values *for all data samples* in the training dataset

	Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
Sample 1	2104	5	1	45	460
Sample 2	1416	3	2	40	232
Sample 3	1534	3	2	30	315
Sample 4	852	2	1	36	178

Outline

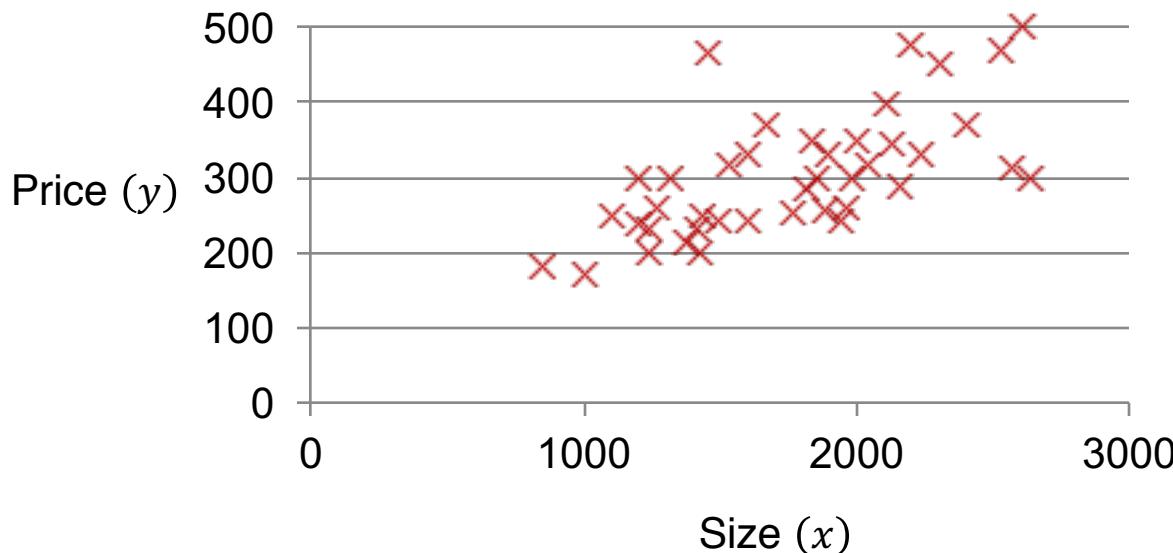
- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization

Model

- For simplicity, first consider only one feature, e.g., *house size*

Size (feet) x	Price (\$ 1000) y
2104	460
1416	232
852	178
....

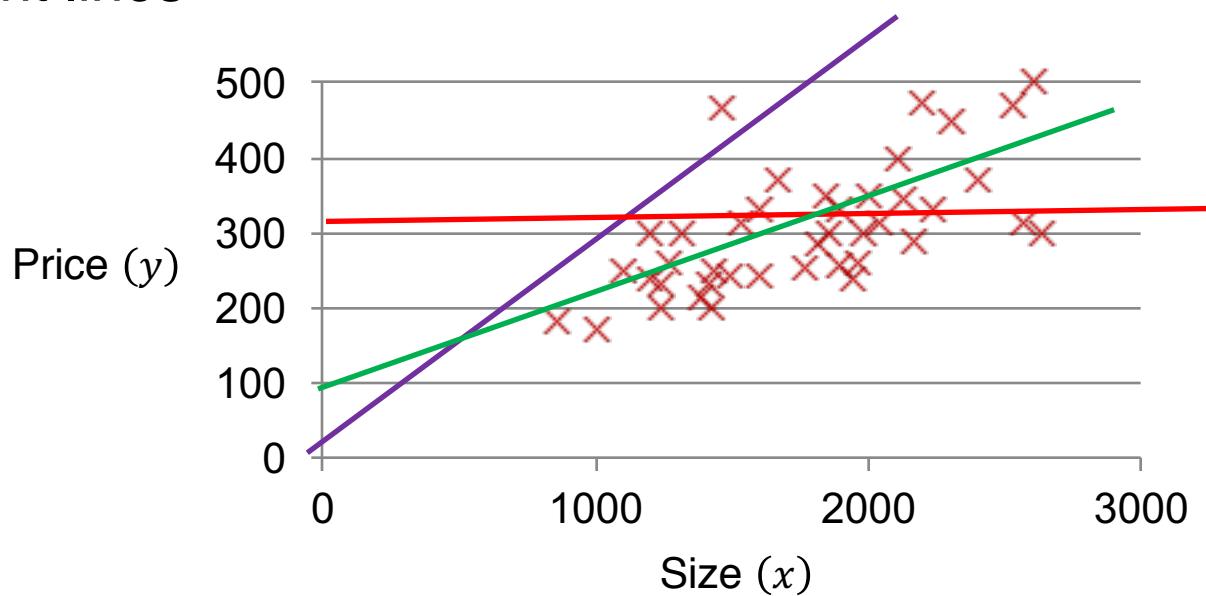
- Plot the (x, y) pairs on a plane



- The prediction function reduces to

$$f(x) = w_0 + w_1 x$$

- For different w_0 and w_1 , the function $f(x)$ represents different straight lines



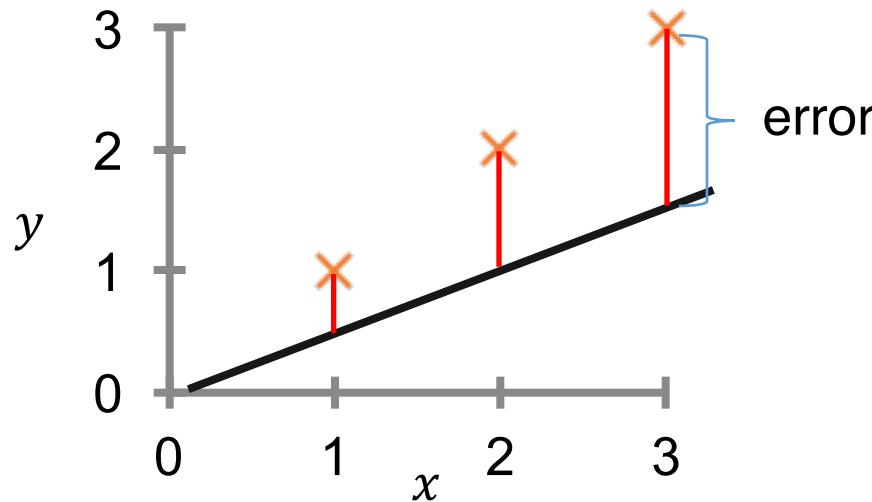
- The objective is to find an appropriate w_0 and w_1 so that **the line is as fit as possible** with the true y 's of all given x

Cost / Loss Function

- Mathematically, the objective can be described as minimizing the cost (loss) function

$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (f(x^{(i)}) - y^{(i)})^2$$

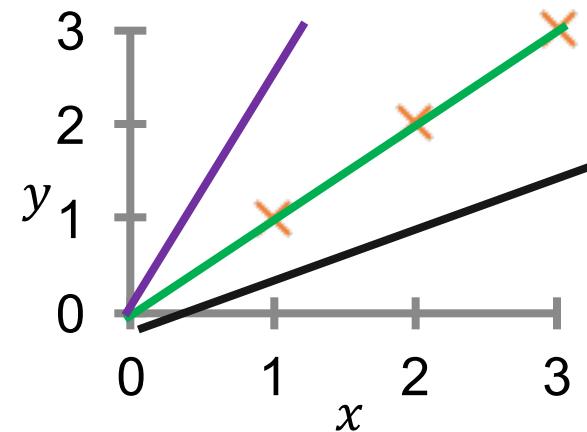
where $x^{(i)}$ and $y^{(i)}$ means the i -th feature and target values;
 n is the number of training examples



- Substituting $f(x^{(i)}) = w_0 + w_1 x^{(i)}$ into $L(w_0, w_1)$ gives

$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)})^2$$

Remark: To better understand this cost function, we simplify it by setting $w_0 = 0$

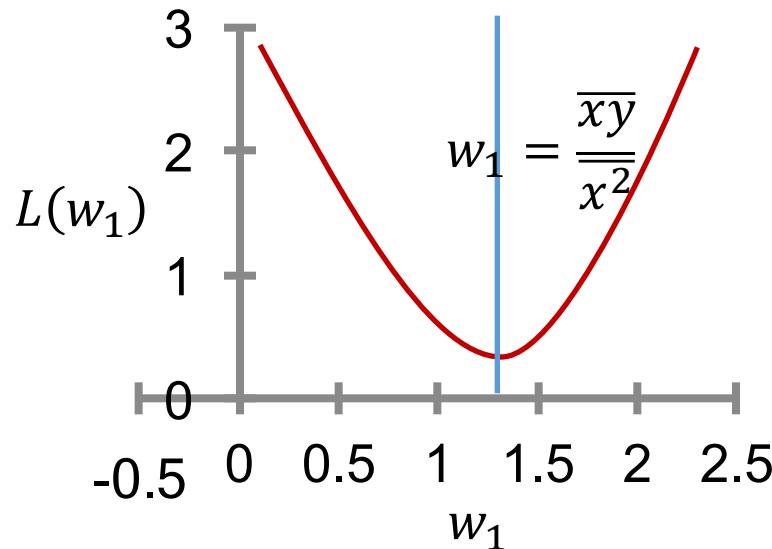


- Then, the cost function becomes

$$L(w_1) = \overline{x^2}w_1^2 - 2\overline{xy}w_1 + \overline{y^2}$$

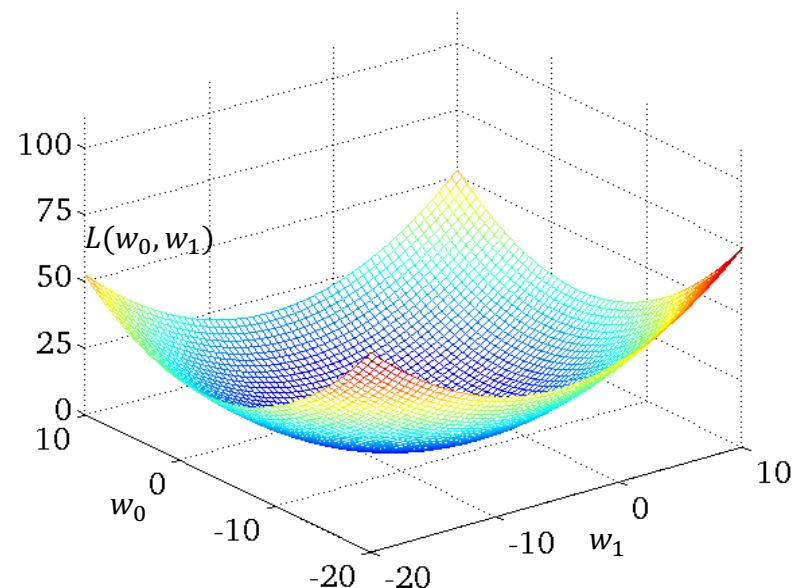
where $\overline{x^2} = \frac{\sum_{i=1}^n (x^{(i)})^2}{n}$, $\overline{xy} = \frac{\sum_{i=1}^n x^{(i)}y^{(i)}}{n}$ and $\overline{y^2} = \frac{\sum_{i=1}^n (y^{(i)})^2}{n}$

- The cost function is a **quadratic function** w.r.t. w_1



$$L(w_1) = \bar{x}^2 w_1^2 - 2\bar{xy}w_1 + \bar{y}^2$$

- If w_0 is taken into account, the cost function $L(w_0, w_1)$ is still a quadratic function, but is **two-dimensional**



$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)})^2$$

- The best w_0 and w_1 can be found by setting the derivatives to zero

$$\frac{\partial L}{\partial w_0} = \frac{2}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)}) = 0$$

$$\frac{\partial L}{\partial w_1} = \frac{2}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)}) x^{(i)} = 0$$

- It can be derived that the best w_0 and w_1 are

$$w_0 = \frac{\bar{xy}\bar{x} - \bar{x}^2\bar{y}}{\bar{x}^2 - \bar{x}^2}$$

$$w_1 = \frac{\bar{x}\bar{y} - \bar{x}\bar{y}}{\bar{x}^2 - \bar{x}^2}$$

Outline

- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization

- Training data from single feature to multiple feature case

Size (feet)	Price (\$ 1000)
x	y

2104	460
------	-----

1416	→ 232
------	-------

1534	315
------	-----

852	178
-----	-----

....
------	------



Size (feet)	# bedrooms	# floors	# years (Ages)	Price (\$ 1000)
x_1	x_2	x_3	x_4	y

2104	5	1	45	460
------	---	---	----	-----

1416	3	2	40	→ 232
------	---	---	----	-------

1534	3	2	30	315
------	---	---	----	-----

852	2	1	36	178
-----	---	---	----	-----

....
------	------	------	------	------

- The function of a general linear regression is

$$f(x_1, x_2 \dots x_m) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_m x_m$$

- x_i is the i -th feature
- Working with the scalar form is cumbersome. Reformulating it into a matrix-form gives

$$f(\mathbf{x}) = \mathbf{x}\mathbf{w}$$

- $\mathbf{x} = [1, x_1, x_2, \dots, x_m]$ is the feature row vector
- $\mathbf{w} = [w_0, w_1, w_2, \dots, w_m]^T$ is the parameter column vector

*By setting the first element in \mathbf{x} to be 1, w_0 can be treated **the same** as the other parameters w_k*

Cost Function

- The objective is still to find a w such that the prediction

$$f(x^{(i)}) = x^{(i)}w$$

is close to the true value $y^{(i)}$, where $x^{(i)}$ and $y^{(i)}$ is the i -th feature vector and target value

Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
....

- Thus, the cost function can be represented as

$$L(w) = \frac{1}{n} \sum_{i=1}^n (x^{(i)}w - y^{(i)})^2$$

- The cost function can be further written as

$$L(\mathbf{w}) = \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

where \mathbf{X} is the feature matrix defined as

$$\mathbf{X} \triangleq \begin{bmatrix} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(n)} \end{bmatrix}$$

	Size (feet)	# bedrooms	# floors	# years (Ages)	Price (\$ 1000)	
	x_0	x_1	x_2	x_3	x_4	y
$\mathbf{X} =$	1	2104	5	1	45	460
	1	1416	3	2	40	232
	1	1534	3	2	30	315
	1	852	2	1	36	178
	1

- The gradient of the cost function w.r.t. \mathbf{w} is

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{2}{n} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

- Since $L(\mathbf{w})$ is a convex function, its optima can be found by setting

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{2}{n} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

- Solving the equation gives

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- It can be verified that when the number of feature is 1, the result reduces to

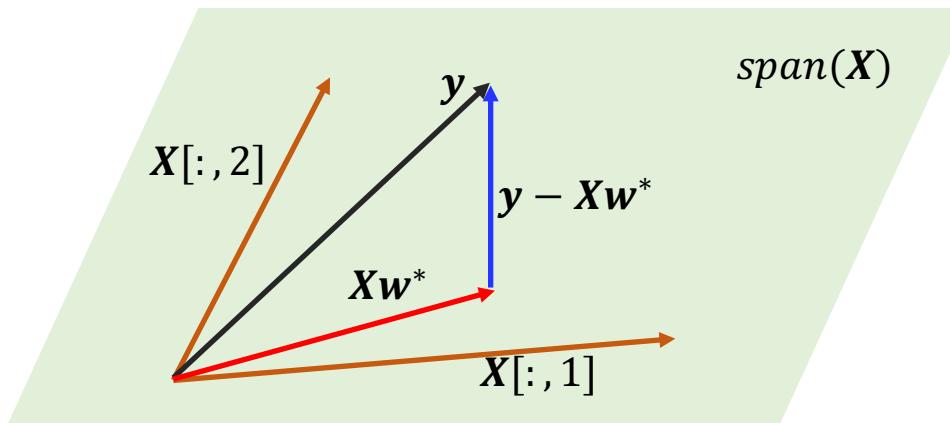
$$w_0 = \frac{\bar{xy} - \bar{x}^2 \bar{y}}{\bar{x}^2 - \bar{x}^2}, \quad w_1 = \frac{\bar{x}\bar{y} - \bar{xy}}{\bar{x}^2 - \bar{x}^2}$$

Geometric Interpretation

- From the requirement of $\mathbf{X}^T(\mathbf{X}\mathbf{w}^* - \mathbf{y}) = \mathbf{0}$, we can see that

$$\mathbf{y} - \mathbf{X}\mathbf{w}^* \perp \text{span}(\mathbf{X})$$

- The result suggests that $\mathbf{X}\mathbf{w}^*$ can be understood as *the projection of \mathbf{y} onto the space spanned by \mathbf{X}*



Outline

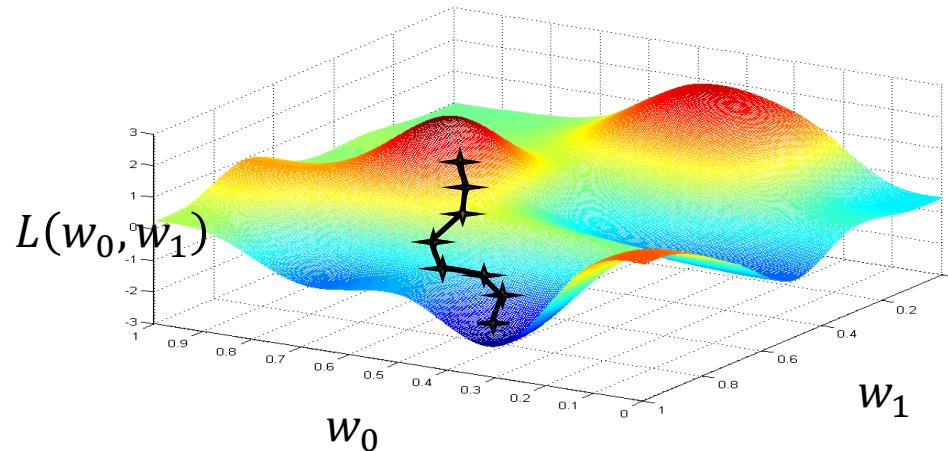
- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization

Gradient Descent

- Analytical solutions do *not always exist*, or evaluating the analytical expression is *computational expensive*
- Then, we should resort to numerical methods, e.g., the gradient descent

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - r \cdot \left. \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^{(t)}}$$

– r : the learning rate

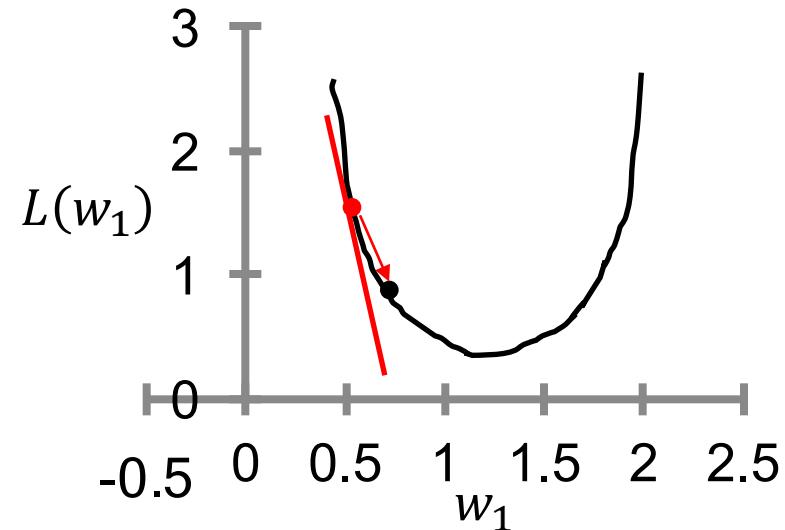
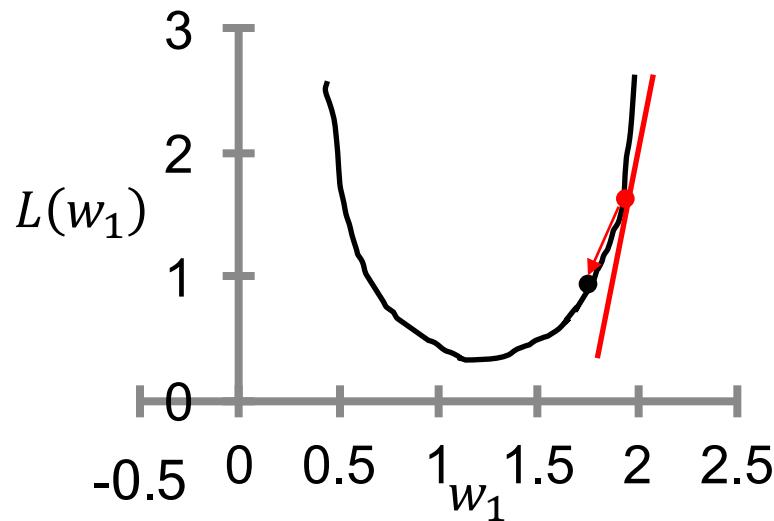


- Let us take the single-feature case and set $w_0 = 0$ as an example, in which the loss function becomes

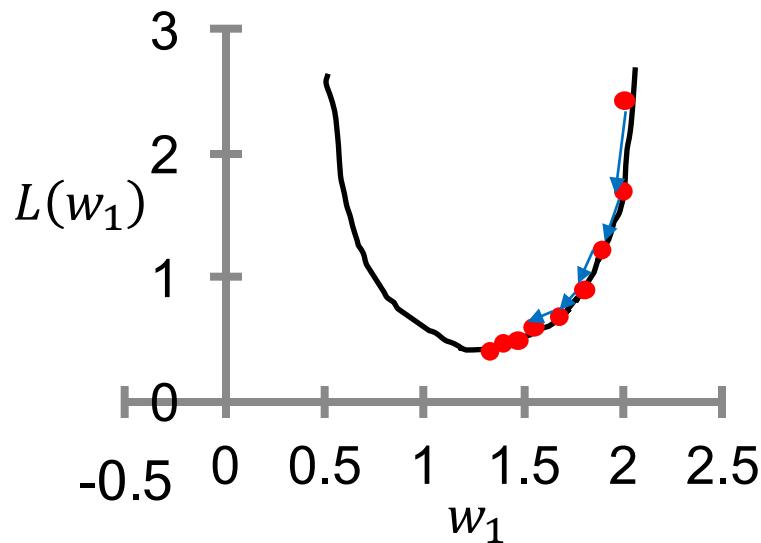
$$L(w_1) = \frac{1}{n} \sum_{i=1}^n (w_1 x^{(i)} - y^{(i)})^2$$

- The parameter w_1 can be updated as

$$w_1^{(t+1)} = w_1^{(t)} - r \cdot \left. \frac{\partial L(w_1)}{\partial w_1} \right|_{w_1=w_1^{(t)}}$$

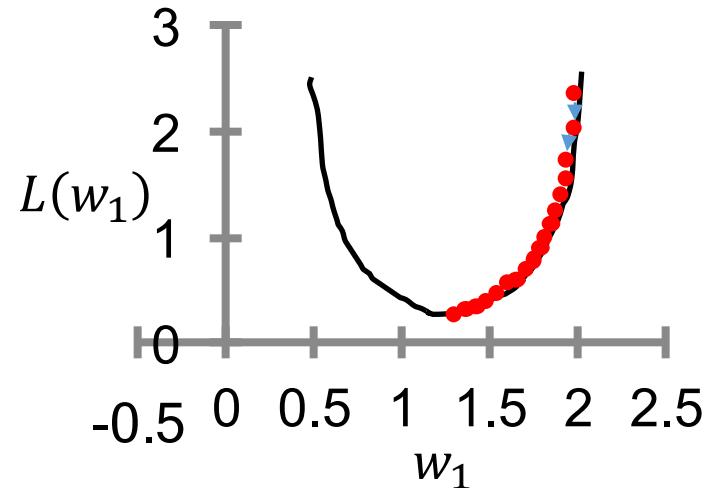


- With appropriate learning rate, the model parameter is iteratively updated, and will eventually converge to the optima

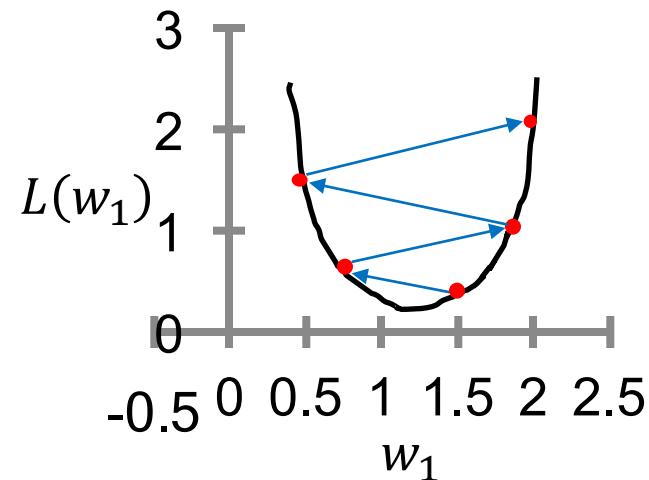


- As approaching the optimal value, the gradient becomes smaller and smaller. Thus, *even if the learning rate is fixed*, the updating intervals also approach 0 as the iteration proceeds, as long as the rate is set appropriately

- If the learning rate is **too small**, the convergence speed will be **very slow**



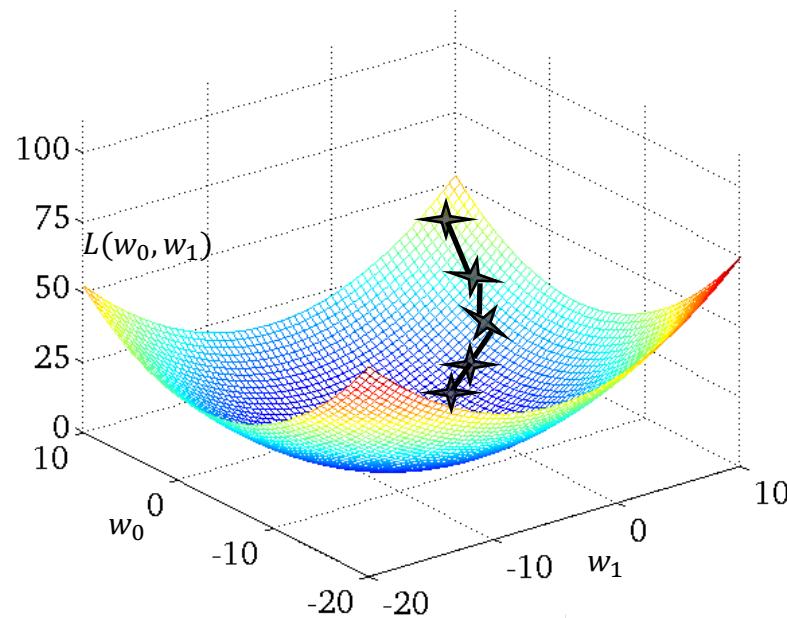
- If the learning rate is **too large**, the iteration may **diverge**
- So, setting appropriate learning rate is important



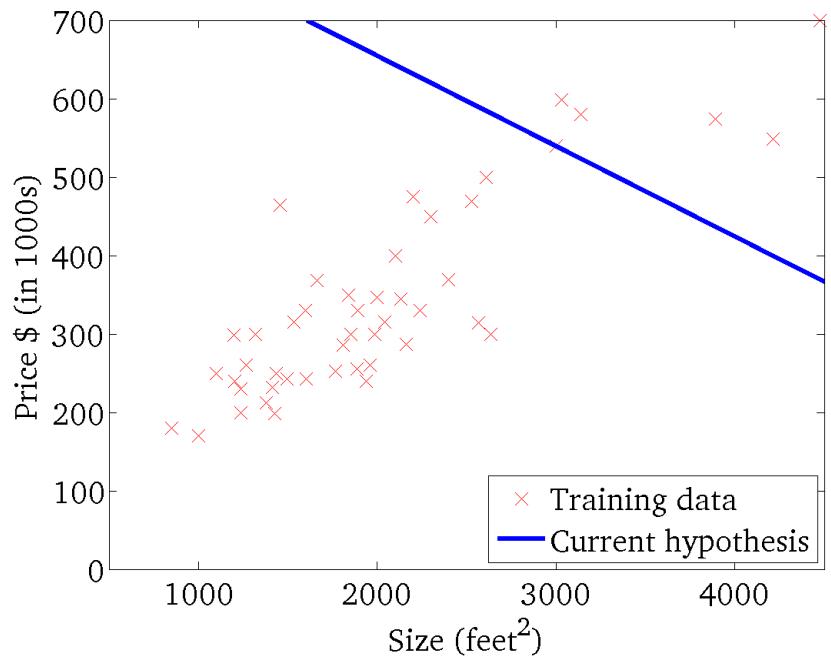
- Now consider the case with both w_0 and w_1

$$w_0^{(t+1)} = w_0^{(t)} - \cancel{r} \cdot \frac{\partial L(w_0, w_1)}{\partial w_0} \Big|_{w_0=w_0^{(t)}, w_1=w_1^{(t)}}$$

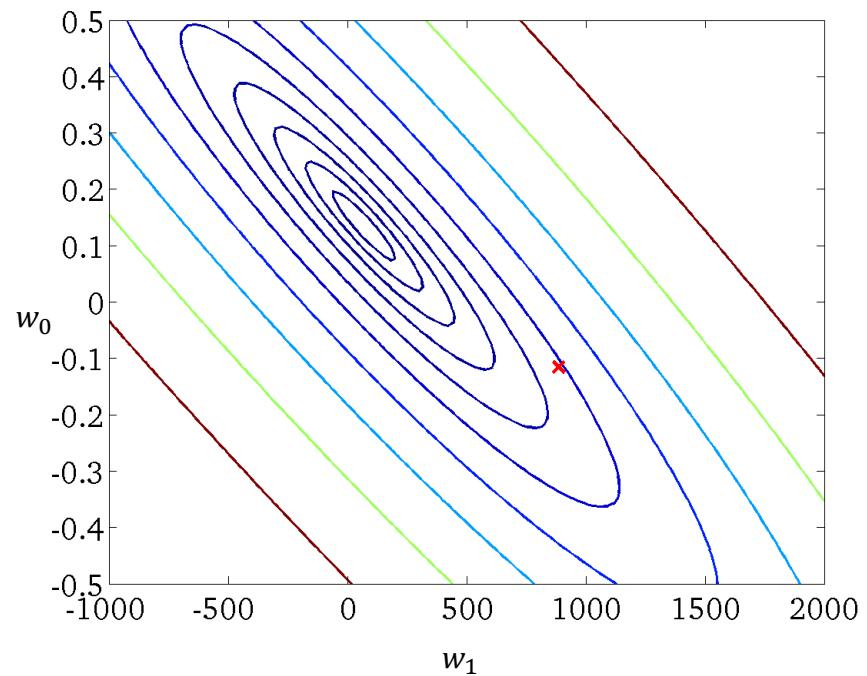
$$w_1^{(t+1)} = w_1^{(t)} - \cancel{r} \cdot \frac{\partial L(w_0, w_1)}{\partial w_1} \Big|_{w_0=w_0^{(t)}, w_1=w_1^{(t)}}$$



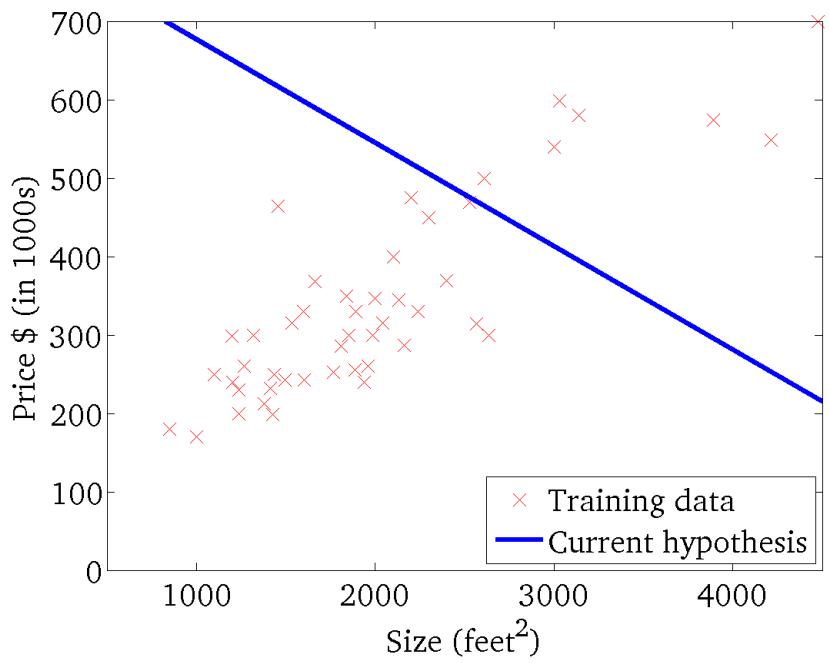
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



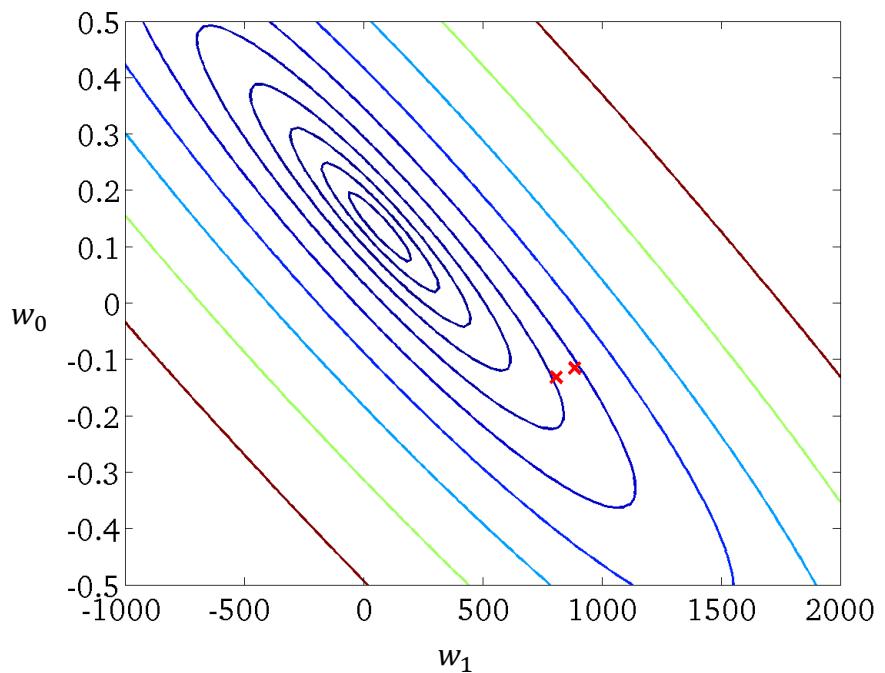
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



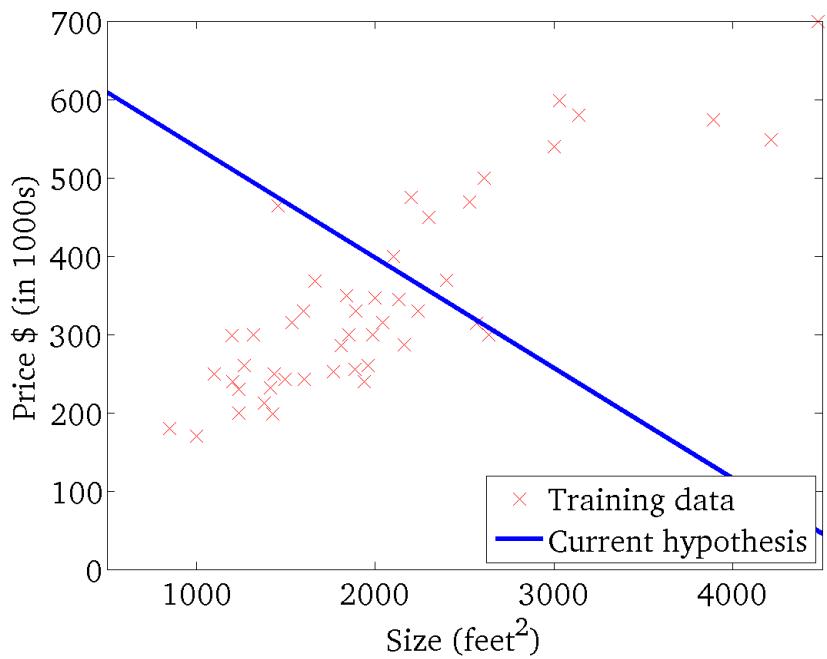
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



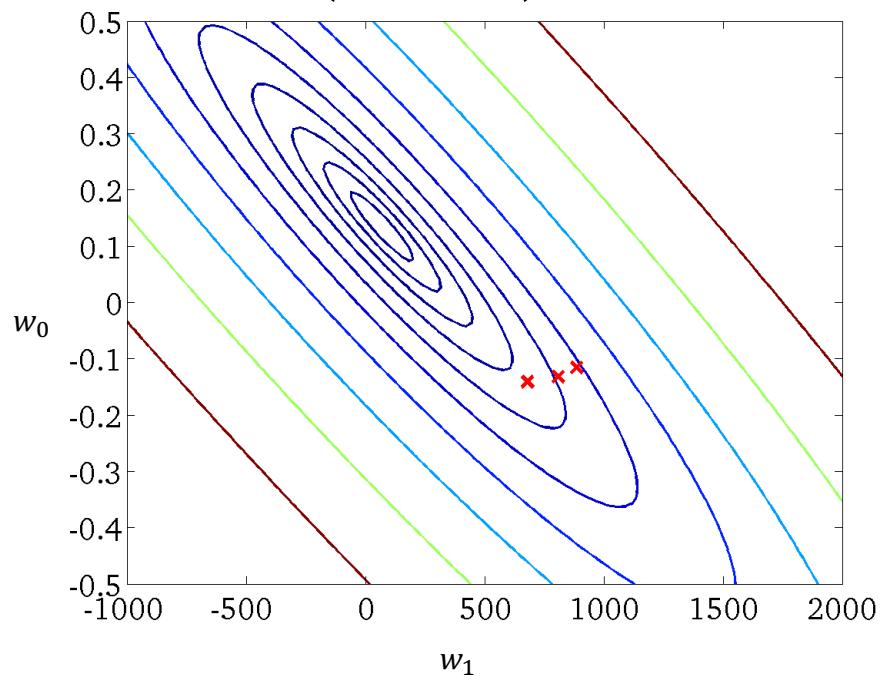
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



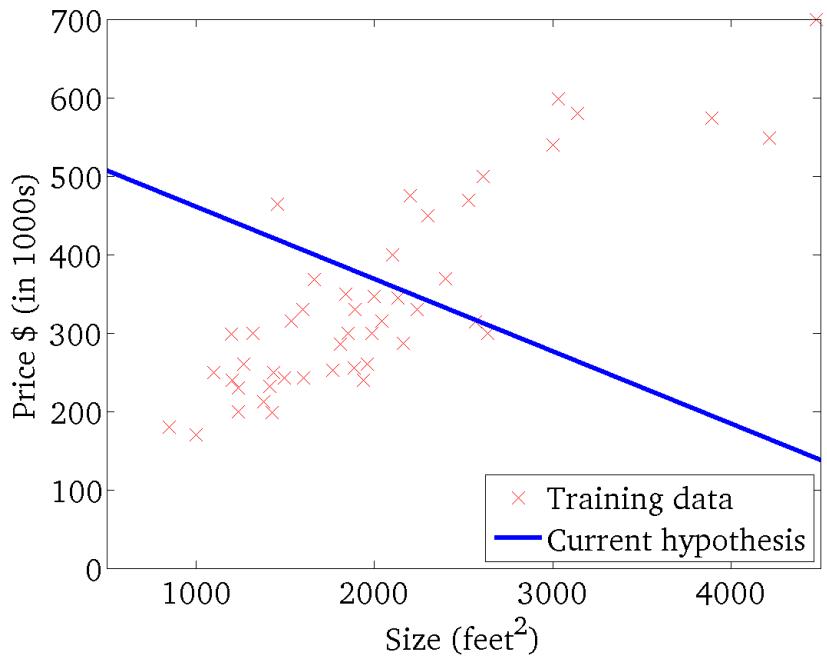
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



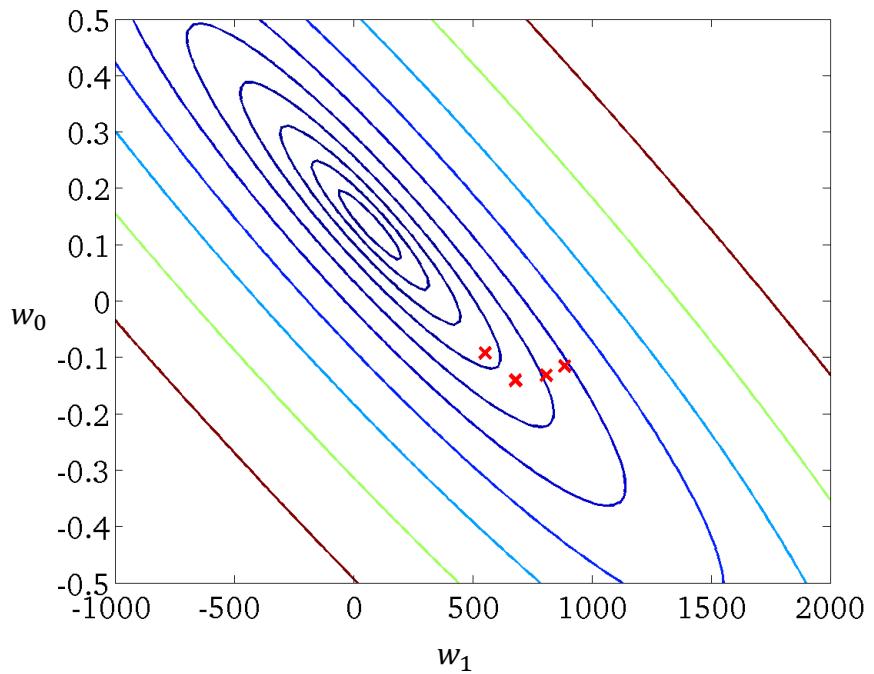
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



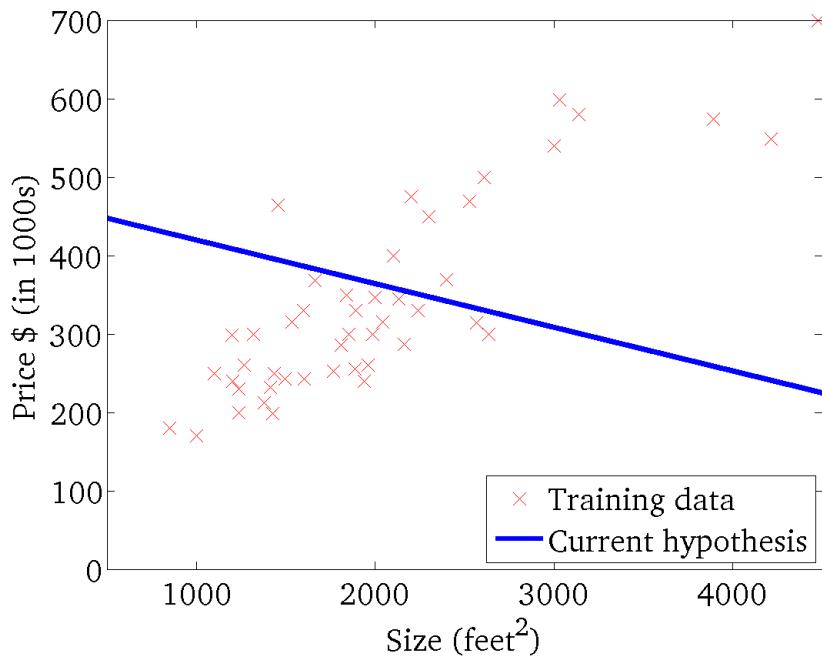
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



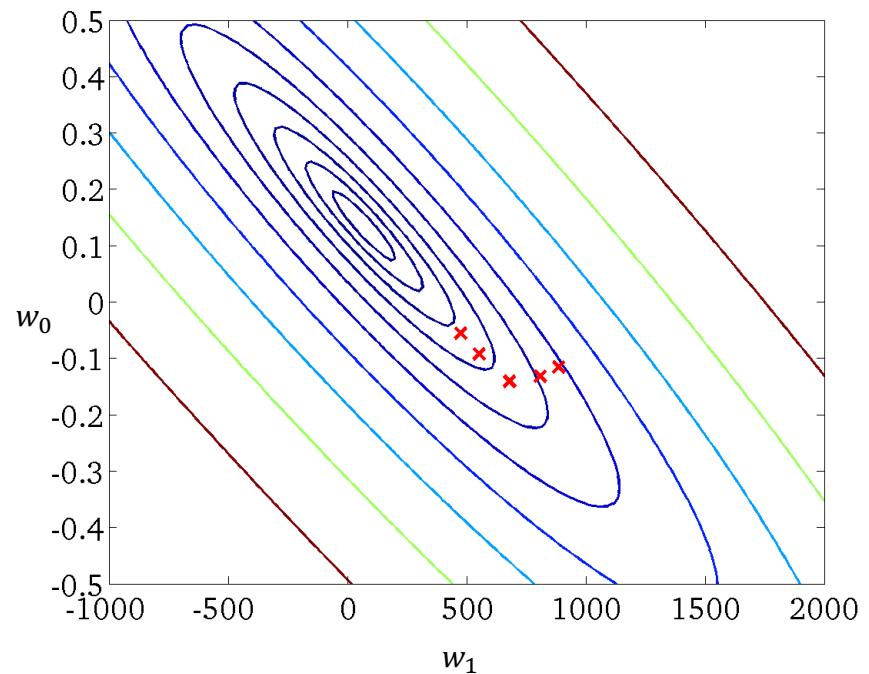
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



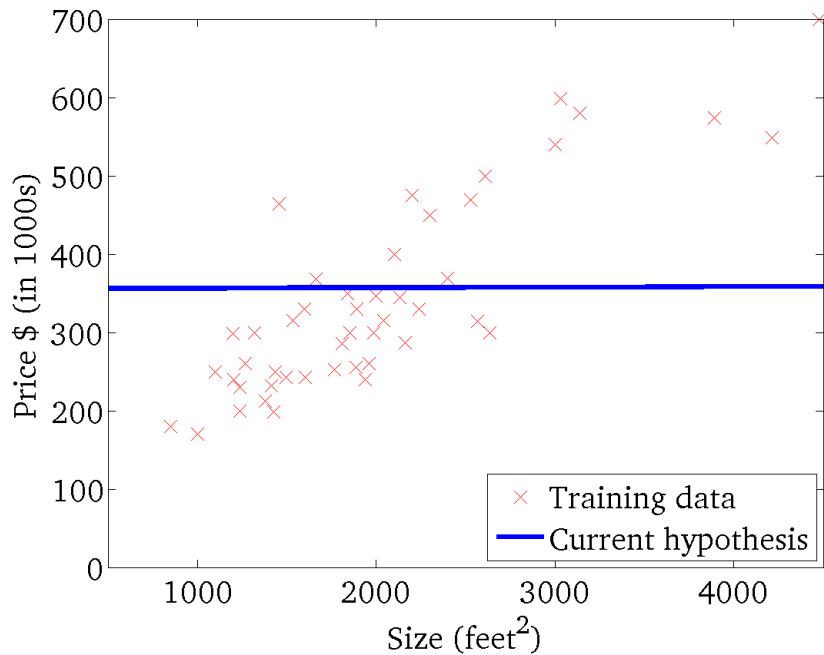
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



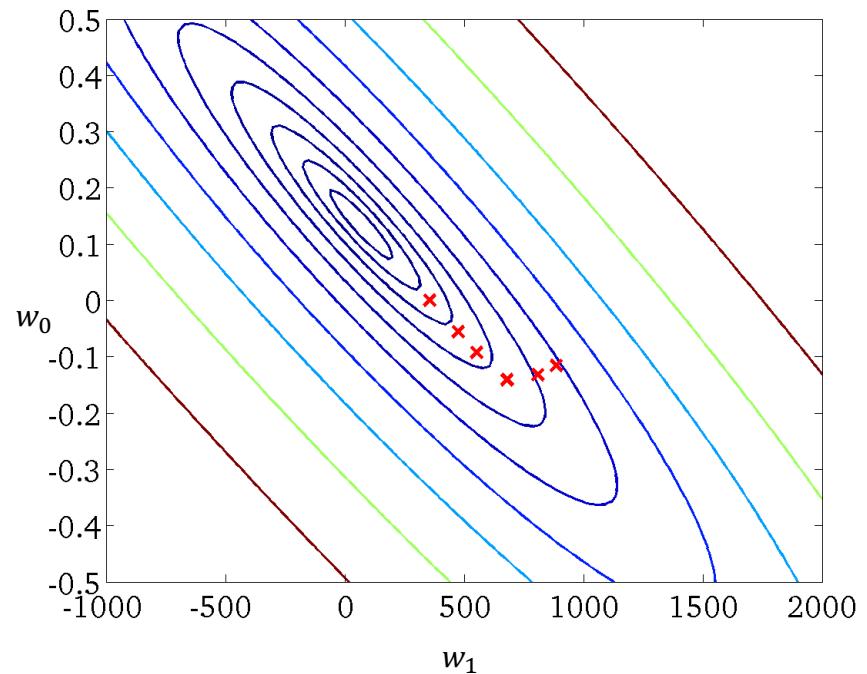
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



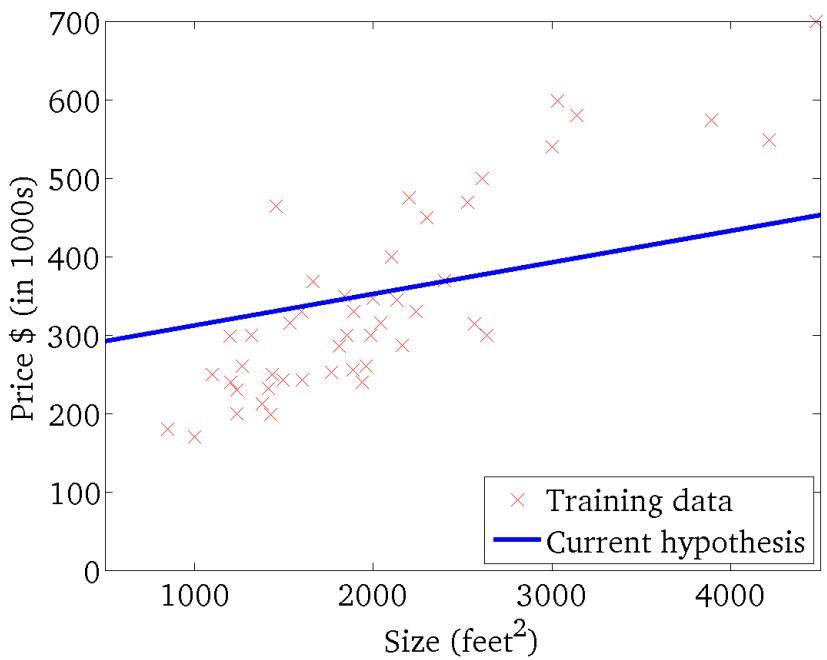
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



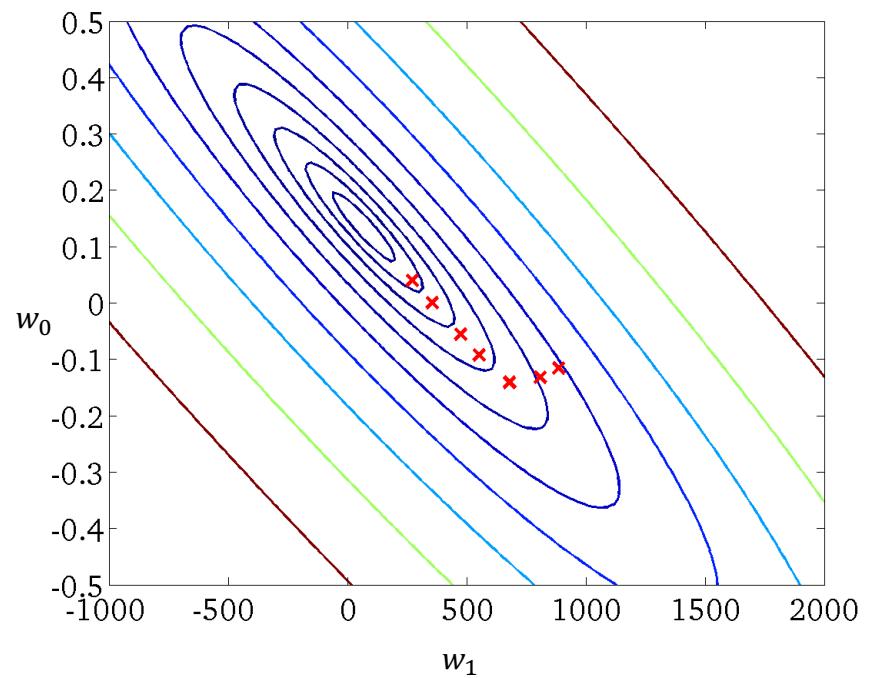
The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$



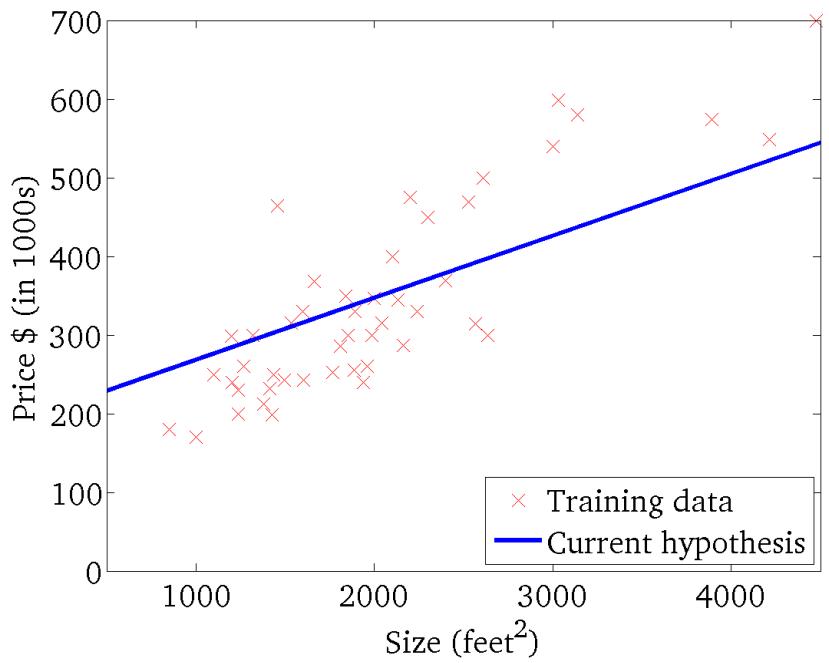
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



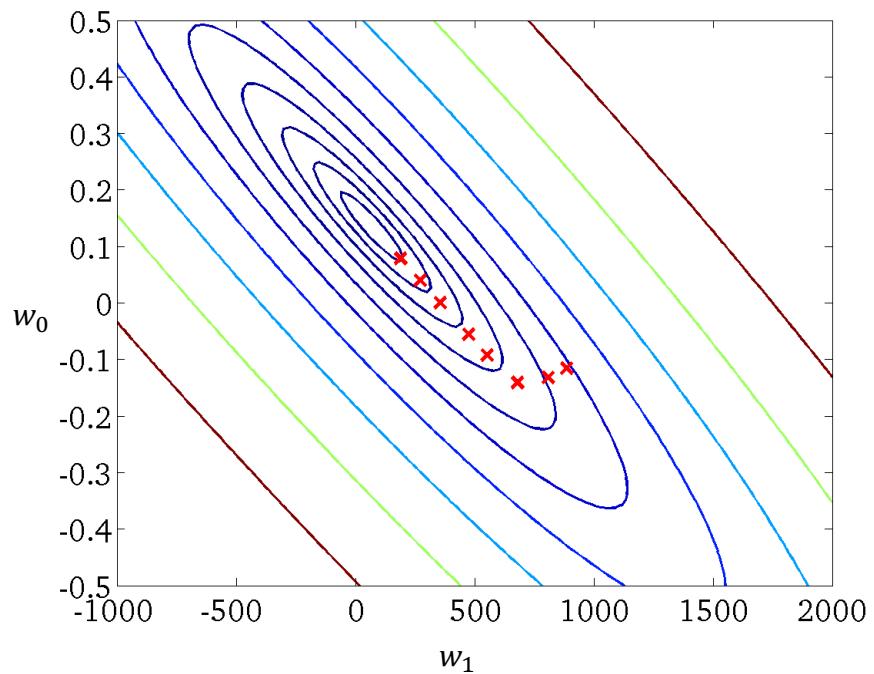
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



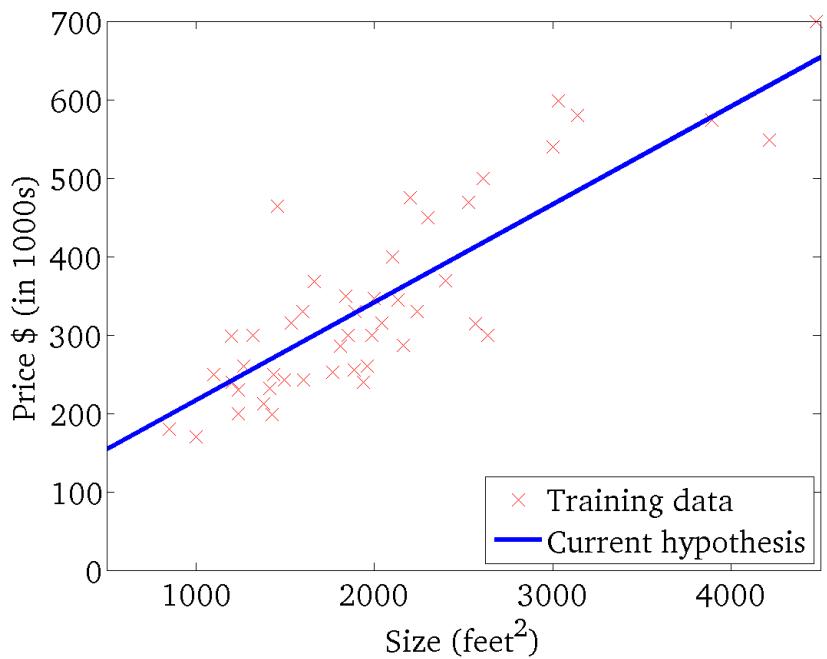
The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



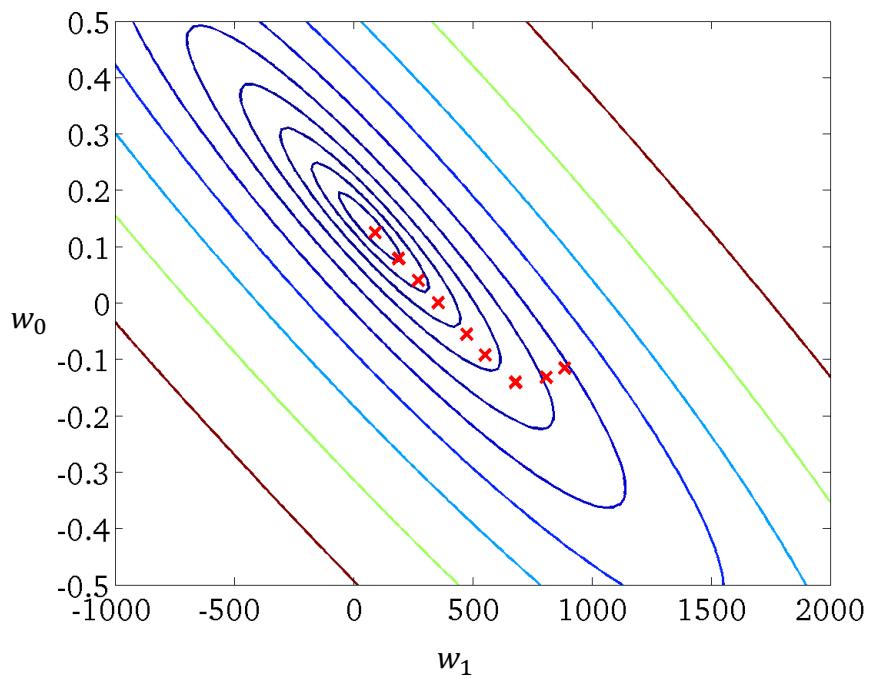
The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)})$



Stochastic Gradient Descent

- The GD algorithm need to evaluate the gradient of loss w.r.t. model parameters w *at every iteration*
- Generally, the gradient takes the form

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{n} \sum_{i=1}^n \frac{\partial \ell(\mathbf{w}, \mathbf{x}^{(i)}, y^{(i)})}{\partial \mathbf{w}}$$

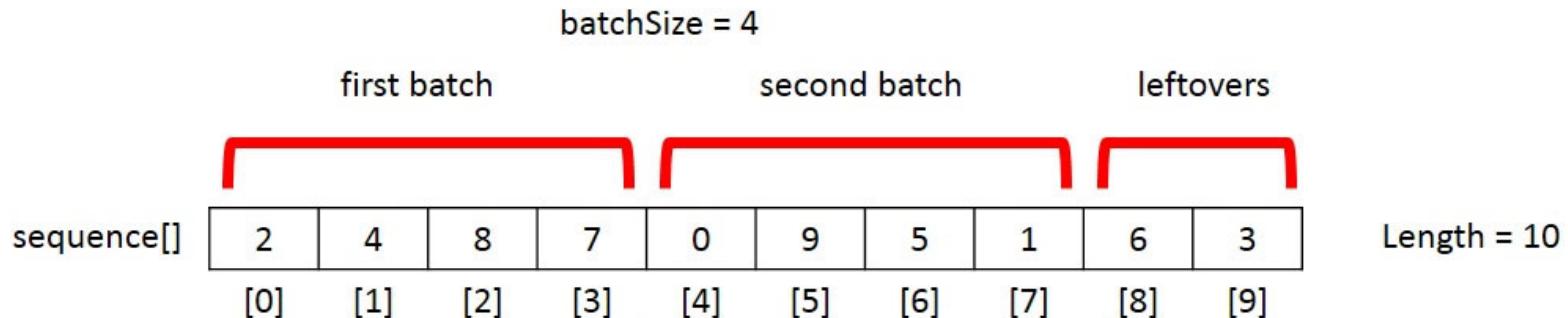
- Every iteration requires computing the gradient for *all data samples in the training dataset*

The complexity would be extremely high for large datasets

- To reduce the complexity, we can estimate the gradient $\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}}$ using a small portion of the dataset, *i.e.* *mini-batch*

- How to obtain the mini-batches?

- Reshuffling
- Segmenting



- Update:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + r \cdot \frac{1}{|\mathcal{B}_t|} \sum_{i \in \mathcal{B}_t} \frac{\partial \ell(\mathbf{w}, \mathbf{x}^{(i)}, y^{(i)})}{\partial \mathbf{w}}$$

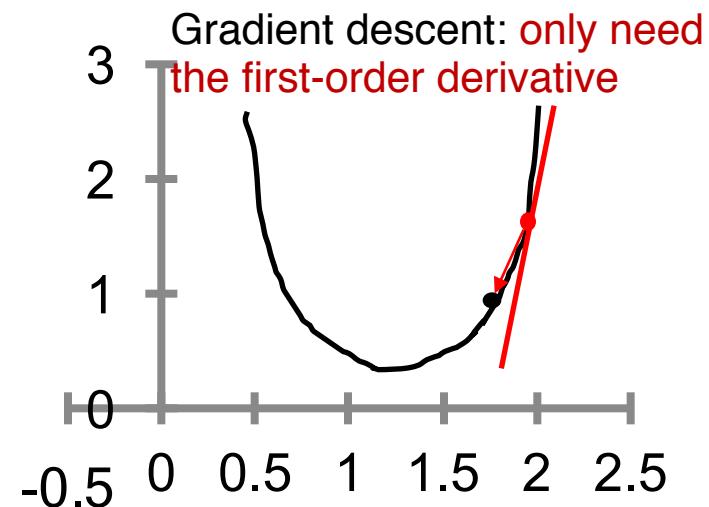
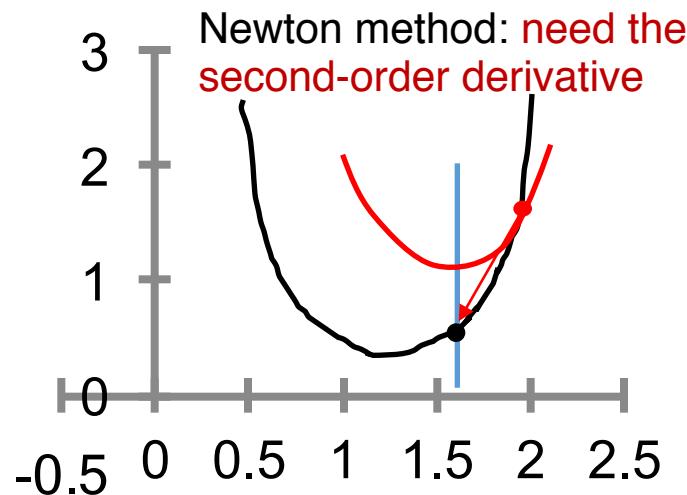
A noisy estimate to
the true gradient

where \mathcal{B}_t is a mini-batch of the dataset at the t -th iteration

Other Optimization Methods

- There also exist many other optimization methods

1) Newton method



Advantages

- No need to manually choose the learning rate
- Faster convergence rate

Disadvantages

- More expensive

- 2) Quasi-Newton methods
- 3) Conjugate gradient method
- 4) Coordinated descent method

⋮

These methods generally converge faster than the gradient method, but are more computationally expensive